

# RUTHVIK VISHWANATHWAR

Senior Data Engineer | Azure Data Engineer | Databricks Engineer  
ruthvikvishwa@gmail.com • 919-397-3265 • [www.linkedin.com/in/ruthvik-vishwa](https://www.linkedin.com/in/ruthvik-vishwa)

---

## SUMMARY

Data engineering leader with 10+ years building the pipelines, platforms, and architectures that businesses rely on. From modernizing legacy SSIS workflows into scalable Databricks solutions to designing HIPAA-compliant lakehouses that process millions of healthcare records daily, the work is always about making data faster, cleaner, and more valuable. Deep hands-on expertise across the Azure ecosystem, Delta Lake, PySpark, and ADF, with a consistent record of turning complex data challenges into production-grade results.

---

## CERTIFICATIONS

- Azure Databricks Platform Architect, Academy Accreditation
- SQL for Data Science

---

## CORE COMPETENCIES

- **Cloud & Platforms:** Azure Data Factory, Databricks, Synapse Analytics, ADLS Gen2, Blob Storage, Logic Apps, DevOps, Microsoft Fabric, OneLake, Snowflake, Power BI
- **Data Engineering:** ETL/ELT Pipelines, Delta Lake, Medallion Architecture, Data Lakehouse, Data Warehousing, SSIS Modernization, Real-Time Streaming (Spark Structured Streaming), Cosmos DB
- **Programming:** Python, PySpark, SQL, T-SQL, DAX, Spark SQL
- **Performance & Optimization:** Z-Ordering, Compaction, Caching, OPTIMIZE/MERGE, Data Skipping, Index Tuning, Query Tuning, Cost Optimization
- **DevOps & Tooling:** Git, GitHub, Azure DevOps CI/CD, GitHub Actions, Databricks Notebooks, JIRA, Confluence, ServiceNow
- **Governance & Compliance:** HIPAA/PHI Masking, Access Controls, Schema Enforcement, Audit Logging, SCDs, Data Quality Validation

---

## PROFESSIONAL EXPERIENCE

**Senior Data Engineer | Databricks Engineer | *Dotcom-Team, Frisco, TX (Remote) (Sentara Health, Georgia Pacific)***  
May 2022 – Present

- Spearheaded end-to-end migration of legacy SSIS packages and SQL Server ETL workflows from Azure Synapse to Azure Databricks, re-engineering over 40+ pipelines into modular PySpark/Spark SQL jobs, cutting average execution time by ~60% and eliminating on-prem infrastructure dependency.
- Architected a multi-source data ingestion framework using ADF and Databricks, ingesting structured and semi-structured data (Parquet, JSON, Avro, XML, Delta) from REST APIs, flat files, SFTP, and third-party systems into ADLS Gen2 at scale, processing millions of records daily.
- Designed and enforced Medallion Architecture (Bronze/Silver/Gold) on Delta Lake, implementing ACID transactions, schema evolution, time travel, and versioning to ensure data reliability and auditability across the entire lakehouse.
- Engineered HIPAA/PHI-compliant data pipelines for large-scale healthcare datasets (claims, member, provider), applying column-level masking, role-based access controls, and PHI tokenization in alignment with enterprise governance standards.
- Drove Delta Lake performance optimization across enterprise datasets using Z-ordering, adaptive partitioning, OPTIMIZE/MERGE, and file compaction, improving query response times by up to 70% on multi-billion row tables.
- Built event-driven and scheduled Databricks Jobs with Delta Live Tables (DLT), automating complex multi-stage data processing workflows with built-in data quality expectations, reducing pipeline failures and manual interventions significantly.
- Reduced a critical batch process from 14 hours to under 3 hours through index tuning, stored procedure refactoring, and query optimization using SQL Server Database Engine Tuning Advisor.
- Developed and maintained CI/CD pipelines in Azure DevOps, automating deployment and version control for ADF, Databricks notebooks, and Power BI semantic models across dev/staging/prod environments.
- Delivered interactive Power BI dashboards and paginated reports for healthcare operations and executive stakeholders, including DAX-driven KPI models and subscription-based delivery.

- Integrated dbt into the Databricks lakehouse to manage and version SQL transformations across Silver and Gold layers, enabling modular, testable data models that accelerated analyst onboarding and reduced ad-hoc query debt.
- Implemented dbt tests (not-null, uniqueness, referential integrity) and source freshness checks across critical data models, embedding data quality gates directly into the transformation pipeline and catching anomalies before they reached downstream reporting layers.

**Stack:** Azure Databricks, PySpark, Spark SQL, ADF, Delta Lake, Delta Live Tables, dbt, Snowflake, ADLS Gen2, SQL Server 2019, Azure Synapse, Power BI, SSIS, Azure DevOps, Python, JIRA

### **Senior Data Engineer** | *E-IT Professionals, Livonia, MI (Blue Cross Blue Shield of Kansas)* May 2021 – May 2022

- Designed and optimized complex SQL Server database objects (stored procedures, views, UDFs) supporting healthcare insurance operations, including member eligibility, claims processing, and provider data workflows, for one of the largest Blue Cross Blue Shield affiliates.
- Replaced manual metric reporting with self-service Power BI dashboards, using Power Query and DAX to model complex insurance KPIs, cutting report turnaround from days to minutes and enabling data-driven decisions at the executive level.
- Established CI/CD pipelines via GitHub Actions for automated testing and deployment of reporting assets, while partnering with business analysts to align insurance KPIs with structured data models, ensuring reliable, production-ready releases tied directly to business outcomes.

**Stack:** SQL Server 2016/2017, Power BI, Power Query, DAX, GitHub Actions, SSMS, JIRA, Confluence

### **Database Developer** | *Skillstorm LLC, Addison, TX (Bank of America)* Nov 2019 – Apr 2021

- Built and deployed SSIS packages for multi-source data extraction and transformation, ingesting from Excel, flat files, and relational databases with Conditional Split, Fuzzy Lookup, and Multicast for data scrubbing and staging validation.
- Designed automated ETL processes ensuring end-to-end data consistency, quality, and scalability across multiple operational systems.
- Authored and maintained SSRS reports with scheduled delivery (daily/weekly/monthly/quarterly), using custom expressions and global variables to satisfy diverse reporting requirements.
- Investigated and resolved performance bottlenecks in stored procedures and complex SQL queries, improving execution plans and reducing load times across key reporting pipelines.

**Stack:** SQL Server 2016/2017, SSIS, SSRS, SSAS, Power BI, Visual Studio 2017, JIRA, Confluence, Bitbucket

### **Database Developer** | *E-IT Professionals Corp, Livonia, MI (Visa Inc.)* Nov 2018 – Oct 2019

- Maintained referential integrity and data consistency using SQL constraints; reviewed and refactored complex multi-join queries to improve performance across high-transaction financial datasets.
- Developed stored procedures, parameterized views, and SSIS packages to support data extraction, blending, and SSRS report generation across multiple source systems.
- Deployed and scheduled periodic SSRS reports with custom expressions, ensuring timely and accurate reporting for operational stakeholders.

**Stack:** SQL Server 2016, SSIS, SSRS, Visual Studio 2017, JIRA, Confluence, Bitbucket

### **Data Engineer** | *Vintech, Missouri* Aug 2016 – Nov 2018

- Delivered Power BI dashboards and SSRS reports with custom DAX measures for senior management, enabling data-driven decision making across business units.
- Authored SQL data manipulation queries and stored procedures aligned to specification, supporting downstream reporting and analytics pipelines.
- Participated in all SDLC phases from requirements analysis through UAT and production deployment in an Agile environment.

**Stack:** Power BI, SQL Server, SSIS, SSRS, DAX, Visual Studio 2017, Git, JIRA

## **EDUCATION**

**M.S. Information Technology** | GPA: 3.8/4.0 | Campbellsville University, Kentucky

**M.S. Computer Science** | GPA: 3.7/4.0 | New York Institute of Technology, Manhattan, NY

**B.S. Information Technology** | GPA: 3.5/4.0 | Scient Institute of Technology, India